llm-d on Intel XPU/GPU: Kubernetes Deployment Whitepaper

Contents

Intel XPU Deployment Whitepaper: llm-d

Target hardware: Intel Data Center GPU (Max series, Arc Pro B60) exposed to Kubernetes via the `gpu.intel.com` Dynamic Resource Allocation (DRA) device class

Table of Contents

- appendix-a-env-vars.md — Appendix A. Environment Variable Reference
- appendix-b-known-issues.md — Appendix B. Known Issues

  `tracking.md` in this directory is a separate, continuously-updated per-case validation-status tracker (hands-on-verified vs. documented-but-untested) — useful when browsing the repository directly, but it is not one of the chapters above and is not included in the compiled PDF/Word export of this whitepaper.

Scope note: which guides actually ship an Intel XPU configuration

A repo-wide search of `guides/**/*.md` and `guides/**/*.yaml` for `xpu`/`intel` found **dedicated Kustomize overlays** (`modelserver/xpu/...`) in exactly six guides: `optimized-baseline`, `pd-disaggregation`, `precise-prefix-cache-routing`, `tiered-prefix-cache`, `wide-ep-lws`, and `multimodal-serving/aggregation`. These are Chapter 1.

The remaining guides (`flow-control`, `predicted-latency-routing`, `workload-autoscaling`, `rollouts`, `agentic-serving`, `multi-model-routing`, `asynchronous-processing`, `batch-gateway`, `no-kubernetes-deployment`) are **router-layer or composite features that are accelerator-agnostic by design** — they sit on top of whatever model-server overlay you deploy (including the Intel XPU overlay from `optimized-baseline`), and reference it rather than shipping a separate XPU overlay of their own (Chapters 2–5). `no-kubernetes-deployment` is explicitly documented as NVIDIA + vLLM specific, with the Intel XPU substitution spelled out in full in its chapter — it is not a maintained, first-class Intel XPU path upstream.

`multimodal-serving/e-disaggregation` (Encode Disaggregation) was investigated and excluded from this whitepaper: its "Supported Hardware Backends" table lists NVIDIA GPU only, with no Intel XPU backend or accelerator-agnostic substitution path — see `tracking.md` for the one-line record. All case-by-case verification status is tracked in `tracking.md`.

Chapter 0. Common Installation

Applies to every case in this whitepaper. Later chapters link back here instead of repeating these steps.

0.1 Cluster and Hardware Requirements

- A Kubernetes cluster (1.29+ recommended; 1.34+ if you plan to use the DRA-based deployment templates, since DRA's structured-parameters API only reached GA in that release) with nodes exposing Intel Data Center GPU Max series or Intel Arc Pro GPUs, and the host Linux Intel GPU kernel driver already installed on each node
- The **Intel GPU DRA driver** installed (see 0.1.1 below), so that `kubectl get deviceclass` lists `gpu.intel.com`
- Local client tools: `kubectl`, `helm` (v3); see `helpers/client-setup/README.md`

0.1.1 Install the Intel GPU DRA Driver

Not part of the llm-d repo itself — this is a cluster-level, vendor-provided component from `intel/intel-resource-drivers-for-kubernetes` that must be installed once per cluster before any Intel XPU case in this whitepaper. Install via its published Helm chart:

```
helm install \
  --namespace intel-gpu-resource-driver \
  --create-namespace \
  intel-gpu-resource-driver oci://ghcr.io/intel/intel-resource-drivers-for-
↪  kubernetes/intel-gpu-resource-driver-chart
```

Verify the driver is up and the `gpu.intel.com` `DeviceClass` and per-node `ResourceSlice` objects were published:

```
kubectl get pods -n intel-gpu-resource-driver
kubectl get deviceclass gpu.intel.com
kubectl get resourceslices
```

See the driver's own GPU documentation for container-runtime CDI requirements (CRI-O 1.23+ or Containerd 1.7+ with CDI enabled) and alternative install methods (kustomize, Node Feature Discovery-scoped rollout).

0.2 Clone the Repository

```
export branch="main"
git clone https://github.com/llm-d/llm-d.git && cd llm-d && git checkout
 ↪  ${branch}
export REPO_ROOT=$(realpath $(git rev-parse --show-toplevel))
source ${REPO_ROOT}/guides/env.sh
```

`guides/env.sh` centralizes the chart versions and OCI addresses shared by every guide (`GAIE_VERSION`, `ROUTER_CHART_VERSION`, `ROUTER_STANDALONE_CHART`, etc.), so commands never hardcode a version.

0.3 Install the Gateway API Inference Extension CRDs

> [!IMPORTANT] CRD installation is a single command below. Gateway Provider installation (Istio / GKE / Agentgateway) lives under `docs/infrastructure/gateway/`.

```
kubectl apply -f https://github.com/kubernetes-sigs/gateway-api-inference-
 ↪  extension/releases/download/${GAIE_VERSION}/v1-manifests.yaml
```

If you plan to use a Kubernetes Gateway (instead of the Router's standalone mode), follow the provider-specific guide under `${REPO_ROOT}/docs/infrastructure/gateway/` once. A single Gateway can be shared across every case in this whitepaper via a dedicated HTTPRoute per case.

0.4 Namespace and HuggingFace Token

```
export NAMESPACE="llm-d-<case-name>"
kubectl create namespace ${NAMESPACE} --dry-run=client -o yaml | kubectl apply
 ↪  -f -

export HF_TOKEN=<your HuggingFace token>
kubectl create secret generic llm-d-hf-token \
    --from-literal="HF_TOKEN=${HF_TOKEN}" \
    --namespace "${NAMESPACE}" \
    --dry-run=client -o yaml | kubectl apply -f -
```

0.5 Proxy Configuration (if required)

If the cluster needs a corporate proxy to reach the HuggingFace Hub, add `HTTP_PROXY` / `HTTPS_PROXY` / `NO_PROXY` (and lowercase equivalents) to the relevant container env. This is applied via a Kustomize patch in the overlay-based cases in this whitepaper, or via `--set` / a values file in Helm-based steps.

1.1 Optimized Baseline

Overview

Deploys the recommended out-of-the-box routing configuration for vLLM/SGLang/TensorRT-LLM: prefix-cache-aware scoring (via the `prefix-cache-scorer` and `no-hit-lru-scorer`) combined with load-aware scoring (`kv-cache-utilization` and `queue-size` scorers). This is the baseline that several other guides (Flow Control, Predicted Latency Routing, Agentic Serving) build on top of. This guide has a dedicated **Intel XPU E2E CI workflow** (`consolidate-status-optimized-baseline- intel-acc-xpu-vllm-x.yaml`), making it the most CI-validated Intel XPU path in the repo today.

Default model: `Qwen/Qwen3-32B` (GPU); the Intel XPU overlay uses a CI-sized `Qwen/Qwen3-0.6B` for fast validation — swap in a production model before real deployment.

Prerequisites

Chapter 0 only. No case-specific prerequisites.

Deployment Steps

Step 1 — Deploy the llm-d Router (Standalone mode)

```
export GUIDE_NAME="optimized-baseline"
export NAMESPACE="llm-d-optimized-baseline"

helm install ${GUIDE_NAME} \
    ${ROUTER_STANDALONE_CHART} \
    -f ${REPO_ROOT}/guides/recipes/router/base.values.yaml \
    -f ${REPO_ROOT}/guides/${GUIDE_NAME}/router/${GUIDE_NAME}.values.yaml \
    -n ${NAMESPACE} --version ${ROUTER_CHART_VERSION}
```

Step 2 — Deploy the Model Server (Intel XPU overlay)

> [!IMPORTANT] The Intel XPU path uses Kubernetes **DRA**: GPUs are declared via a ResourceClaimTemplate (`deviceClassName: gpu.intel.com`), not `resources.limits`. Confirm the cluster has DRA enabled and the `gpu.intel.com` device class available before deploying.

```
export ACCELERATOR_TYPE=xpu
export MODEL_SERVER=vllm
kubectl apply -n ${NAMESPACE} -k ${REPO_ROOT}/guides/${GUIDE_NAME}
 ↪ /modelserver/${ACCELERATOR_TYPE}/${MODEL_SERVER}/
```

> [!NOTE] Unlike the NVIDIA GPU path, the Intel XPU overlay path has **no** `${INFRA_PROVIDER}` suffix — apply the directory directly.

This overlay (`modelserver/xpu/vllm/`): - uses `recipes/modelserver/base/single-host/default` as its base and the `recipes/modelserver/components/images/xpu-vllm` image component - patches the decode Deployment (`patch-vllm.yaml`) to run `vllm serve Qwen/Qwen3-0.6B --dtype=float16 --disable-sliding-window ...` - adds a ResourceClaimTemplate (`resource-claim-template.yaml`) requesting 1 `gpu.intel.com` device per replica, with `fsGroup: 107` / `supplementalGroups: [107]` for Intel GPU device node access

Step 3 — (Optional) Enable Monitoring

```
kubectl apply -n ${NAMESPACE} -k
 ↪ ${REPO_ROOT}/guides/recipes/modelserver/components/monitoring
```

Verification

```
export IP=$(kubectl get service ${GUIDE_NAME}-epp -n ${NAMESPACE} -o
 ↪  jsonpath='{.spec.clusterIP}')
kubectl get pods -n ${NAMESPACE}
```

Inference Test

```
kubectl run curl-debug --rm -it \
    --image=cfmanteiga/alpine-bash-curl-jq \
    --namespace="${NAMESPACE}" \
    --env="IP=${IP}" \
    -- /bin/bash

curl -X POST http://${IP}/v1/completions \
    -H 'Content-Type: application/json' \
    -d '{"model": "Qwen/Qwen3-0.6B", "prompt": "How are you today?"}' | jq
```

Troubleshooting

- **`ResourceClaim` stuck `Pending`**: confirm the node has registered the `gpu.intel.com` device class; `kubectl describe resourceclaim <name> -n ${NAMESPACE}`.
- **vLLM usage-telemetry write failures under restricted `SecurityContext`** (e.g. OpenShift): set `DO_NOT_TRACK=1` in the container env (already applied in the precise-prefix-cache-routing XPU overlay; port the same fix here if you hit it).

Cleanup

```
helm uninstall ${GUIDE_NAME} -n ${NAMESPACE}
kubectl delete -n ${NAMESPACE} -k ${REPO_ROOT}/guides/${GUIDE_NAME}
 ↪  /modelserver/${ACCELERATOR_TYPE}/${MODEL_SERVER}/
kubectl delete namespace ${NAMESPACE}
```

Configuration Reference

Parameter	Value	Purpose
`--dtype`	`float16`	Intel XPU overlay dtype
`--disable-sliding-window`	set	Required by the CI-sized model
`deviceClassName`	`gpu.intel.com`	DRA device class
`fsGroup` / `supplementalGroups`	`107`	Intel GPU render-node group access

Resource requirements (Intel XPU overlay, `Qwen/Qwen3-0.6B`, 2 replicas):

Role	Replicas	CPU (req/limit)	Memory (req/limit)	GPU
Decode	2	4 / 8 cores	12Gi / 24Gi	1× `gpu.intel.com` per replica (DRA claim)

1.2 P/D Disaggregation

Overview

Splits inference into independently-scaled Prefill and Decode deployments, connected via NIXL KV transfer. Native to the llm-d Router, so it composes with prefix-cache-aware and load-aware routing from Optimized Baseline. Default GPU example uses `openai/gpt-oss-120b` (8× TP=1 Prefill + 2× TP=4 Decode); the Intel XPU overlay uses the smaller `Qwen/Qwen3-0.6B` for fast validation.

Best suited for: medium-large models, long input sequences (e.g. 10k ISL / 1k OSL), sparse MoE architectures.

Prerequisites

Chapter 0 only.

Deployment Steps

Step 1 — Deploy the llm-d Router (Standalone mode)

```
export GUIDE_NAME="pd-disaggregation"
export NAMESPACE="llm-d-pd-disaggregation"
export MODEL_NAME="Qwen/Qwen3-0.6B"   # Intel XPU overlay default; swap for
 ↪  production models

helm install ${GUIDE_NAME} \
    ${ROUTER_STANDALONE_CHART} \
    -f ${REPO_ROOT}/guides/recipes/router/base.values.yaml \
    -f ${REPO_ROOT}/guides/${GUIDE_NAME}/router/${GUIDE_NAME}.values.yaml \
    -n ${NAMESPACE} --version ${ROUTER_CHART_VERSION}
```

Step 2 — Deploy the Model Server (Intel XPU overlay)

> [!IMPORTANT] GPUs are declared via DRA `ResourceClaimTemplate` (`deviceClassName: gpu.intel.com`), not `resources.limits."gpu.intel.com/xe"`.

```
kubectl apply -n ${NAMESPACE} -k
 ↪  ${REPO_ROOT}/guides/pd-disaggregation/modelserver/xpu/vllm
```

This overlay includes: - `patch-prefill.yaml` / `patch-decode.yaml`: vLLM launch args, including `--kv-transfer-config` with `kv_connector: NixlConnector`, `kv_role: kv_both`, `kv_buffer_device: cpu`; and env vars `VLLM_USE_V1=1`, `VLLM_NIXL_SIDE_CHANNEL_HOST` (from `status.podIP`), `VLLM_NIXL_SIDE_CHANNEL_PORT=5600`, `UCX_TLS=tcp`, `VLLM_WORKER_MULTIPROC_METHO
- `resource-claim-templates.yaml`: two `ResourceClaimTemplates`, `xpu-prefill-claim` and `xpu-decode-claim`, each requesting 1 `gpu.intel.com` device

For clusters with RDMA connectivity, use `modelserver/xpu/vllm-rdma` instead (UCX transport `ib,rc,ze_copy`).

Step 3 — (Optional) Enable Monitoring

```
kubectl apply -n ${NAMESPACE} -k
 ↪  ${REPO_ROOT}/guides/recipes/modelserver/components/monitoring-pd
```

Verification

```
kubectl get pods -n ${NAMESPACE}
kubectl get pods -n ${NAMESPACE} -l llm-d.ai/role=decode -w
kubectl get pods -n ${NAMESPACE} -l llm-d.ai/role=prefill -w

export IP=$(kubectl get service ${GUIDE_NAME}-epp -n ${NAMESPACE} -o
 ↪  jsonpath='{.spec.clusterIP}')
```

Inference Test

```
kubectl run curl-debug --rm -it \
    --image=cfmanteiga/alpine-bash-curl-jq \
    --namespace="${NAMESPACE}" \
    --env="IP=${IP}" \
    -- /bin/bash

curl -X POST http://${IP}/v1/completions \
    -H 'Content-Type: application/json' \
    -d '{"model": "Qwen/Qwen3-0.6B", "prompt": "How are you today?"}' | jq
```

Troubleshooting

- **Pod stuck at Init:0/1**: known routing-proxy sidecar issue (llm-d/llm-d#241); workaround: set `proxy.enabled: false` — the Gateway API + InferencePool routing already covers what the sidecar provided.
- **KV transfer failures**: check decode/prefill container logs for `nixl`: `kubectl logs -n ${NAMESPACE} <pod> -c modelserver | grep -i nixl`.
- **ResourceClaim stuck Pending**: confirm the `gpu.intel.com` device class is registered; `kubectl get resourceclaims -n ${NAMESPACE}` for details.
- **HuggingFace download failures**: check proxy env vars and DNS (`nslookup hugging-face.co`).

Cleanup

```
helm uninstall ${GUIDE_NAME} -n ${NAMESPACE}
kubectl delete -n ${NAMESPACE} -k
 ↪  ${REPO_ROOT}/guides/pd-disaggregation/modelserver/xpu/vllm
```

Configuration Reference

Parameter	Value	Purpose
`kv_connector`	`NixlConnector`	KV transfer connector
`kv_role`	`kv_both`	Both send and receive

Parameter	Value	Purpose
`kv_buffer_device`	`cpu`	Reduces GPU memory pressure
`VLLM_USE_V1`	`1`	Required for P/D
`VLLM_WORKER_MULTIPROC_METHOD`	`spawn`	Multiprocess start method
`UCX_TLS`	`tcp` (or `ib,rc,ze_copy` for the RDMA overlay)	UCX transport

Resource requirements (Intel XPU overlay, `Qwen/Qwen3-0.6B`):

Role	Replicas	CPU	Memory	GPU
Decode	1	16 cores	64Gi	1× `gpu.intel.com` (DRA claim)
Prefill	3	16 cores/replica	64Gi/replica	1× `gpu.intel.com`/replica (DRA claim)

> [!WARNING] Production models (e.g. `openai/gpt-oss-120b`) require substantially more CPU/memory/GPU than the table above — re-plan capacity for your target model.

1.3 Precise Prefix Cache Routing

Overview

Routes on precise, per-pod KV-cache state instead of traffic heuristics. Each vLLM pod publishes KV-cache events over ZMQ; the router indexes them by block hash and scores candidates by resident prefix fraction, combined with load-aware scoring. Intel XPU overlay uses the CI-sized `Qwen/Qwen3-0.6B` model; update the router's `token-producer modelName` in `router/precise-prefix-cache-routing.values.yaml` to match whatever model you deploy — precise scoring only works if the two match.

Prerequisites

Chapter 0 only. Note `--block-size` in the vLLM args must match the scorer's `tokenProcessorConfig.blockSize` (default 64).

Deployment Steps

Step 1 — Deploy the llm-d Router

```
export GUIDE_NAME="precise-prefix-cache-routing"
export NAMESPACE="llm-d-precise-prefix-cache-routing"

helm install ${GUIDE_NAME} \
    ${ROUTER_STANDALONE_CHART} \
    -f ${REPO_ROOT}/guides/recipes/router/base.values.yaml \
```

```
    -f ${REPO_ROOT}/guides/${GUIDE_NAME}/router/${GUIDE_NAME}.values.yaml \
    -n ${NAMESPACE} --version ${ROUTER_CHART_VERSION}
```

[!NOTE] The router defaults to **active-active HA** (2 EPP replicas), each subscribing to every vLLM pod so both indexes converge independently. Set `--set router.epp.replicas=1` for small fleets or smoke tests.

Step 2 — Deploy the Model Server (Intel XPU overlay)

```
kubectl apply -n ${NAMESPACE} -k
 ↪  ${REPO_ROOT}/guides/precise-prefix-cache-routing/modelserver/xpu/vllm
```

Key args baked into the overlay's `patch-vllm.yaml`:

```
args:
  - "Qwen/Qwen3-0.6B"
  - "--dtype=float16"
  - "--disable-sliding-window"
  - "--block-size=64"
  - "--kv-events-config"
  - |
 ↪  '{"enable_kv_cache_events":true,"publisher":"zmq","endpoint":"$(KV_EVENTS_ENDPOINT)","
 ↪  0.6B"}'
env:
  - name: KV_EVENTS_ENDPOINT
    value: "tcp://*:5556"        # pod-discovery mode: every router replica
 ↪  dials each pod directly
  - name: DO_NOT_TRACK
    value: "1"                   # disables vLLM telemetry writes to /.config,
 ↪  which fail under
                                 # restricted SecurityContexts (e.g. OpenShift)
```

Port 5556 (`kv-events`) is exposed on the pod so every router replica can reach the per-pod ZMQ socket.

Verification

```
kubectl get pods -n ${NAMESPACE}
export IP=$(kubectl get service ${GUIDE_NAME}-epp -n ${NAMESPACE} -o
 ↪  jsonpath='{.spec.clusterIP}')
```

Inference Test

```
kubectl run curl-debug --rm -it \
    --image=cfmanteiga/alpine-bash-curl-jq \
    --namespace="${NAMESPACE}" \
    --env="IP=${IP}" \
    -- /bin/bash

curl -X POST http://${IP}/v1/completions \
    -H 'Content-Type: application/json' \
    -d '{"model": "Qwen/Qwen3-0.6B", "prompt": "How are you today?"}' | jq
```

Troubleshooting

- **Scores look wrong / no cache hits registered**: confirm `router/precise-prefix-cache-routing.values.yaml`'s `modelName` matches the model actually deployed by

the overlay.
- **vLLM telemetry write errors under OpenShift-style restricted `SecurityContext`**: confirm `DO_NOT_TRACK=1` is set.
- **Router not receiving KV events**: verify port 5556 is reachable from the router pod to each vLLM pod (`kubectl exec` into the router pod and test connectivity).

Cleanup

```
helm uninstall ${GUIDE_NAME} -n ${NAMESPACE}
kubectl delete -n ${NAMESPACE} -k
 ↪  ${REPO_ROOT}/guides/${GUIDE_NAME}/modelserver/xpu/vllm
```

Configuration Reference

Parameter	Value	Purpose
`--block-size`	64	Must match scorer `tokenProcessorConfig.blockSize`
`enable_kv_cache_events`	`true`	Turns on KV-cache event publishing
`publisher`	`zmq`	Event transport
`KV_EVENTS_ENDPOINT`	`tcp://*:5556`	Pod-discovery mode ZMQ bind address
`DO_NOT_TRACK`	1	Disables vLLM telemetry (avoids write failures under restricted SCCs)

1.4 Tiered Prefix Cache

Overview

Offloads evicted KV-cache blocks from accelerator memory to larger, cheaper tiers (CPU RAM, and optionally shared filesystem), increasing effective cache size for multi-turn/long-context workloads. Multiple offloading implementations exist (vLLM native `OffloadingConnector`, LMCache, MooncakeStore, SGLang HiCache); the Intel XPU overlay ships the **vLLM native** and **LMCache connector (CPU tier)** paths.

Prerequisites

Chapter 0, plus the prefix-aware routing from Optimized Baseline (already the default scorer stack used here).

Deployment Steps

Step 1 — Deploy the llm-d Router

```
export GUIDE_NAME="tiered-prefix-cache"
export NAMESPACE="llm-d-tiered-prefix-cache"
```

```
helm install ${GUIDE_NAME} \
    ${ROUTER_STANDALONE_CHART} \
    -f ${REPO_ROOT}/guides/recipes/router/base.values.yaml \
    -f ${REPO_ROOT}/guides/${GUIDE_NAME}/router/${GUIDE_NAME}.values.yaml \
    -n ${NAMESPACE} --version ${ROUTER_CHART_VERSION}
```

Step 2 — Deploy the Model Server (Intel XPU overlay, choose a path)

Native (vLLM `OffloadingConnector`, CPU RAM tier):

```
kubectl apply -n ${NAMESPACE} -k
 ↪  ${REPO_ROOT}/guides/tiered-prefix-cache/modelserver/xpu/vllm/base
```

LMCache connector (CPU RAM tier):

```
kubectl apply -n ${NAMESPACE} -k ${REPO_ROOT}/guides/tiered-prefix-
 ↪  cache/modelserver/xpu/vllm/lmcache-connector/cpu/base
```

> [!IMPORTANT] The Intel XPU base overlay sets `securityContext.privileged: true` on the model-server container in addition to the DRA `ResourceClaimTemplate` (`xpu-vllm-gpu`) — this is required for the offloading connector to manage host-pinned memory. Review your cluster's Pod Security Admission policy before applying.

Verification

```
kubectl get pods -n ${NAMESPACE}
export IP=$(kubectl get service ${GUIDE_NAME}-epp -n ${NAMESPACE} -o
 ↪  jsonpath='{.spec.clusterIP}')
```

Inference Test

```
kubectl run curl-debug --rm -it \
    --image=cfmanteiga/alpine-bash-curl-jq \
    --namespace="${NAMESPACE}" \
    --env="IP=${IP}" \
    -- /bin/bash

curl -X POST http://${IP}/v1/completions \
    -H 'Content-Type: application/json' \
    -d '{"model": "Qwen/Qwen3-0.6B", "prompt": "Summarize the benefits of
     ↪  KV-cache offloading."}' | jq
```

Send a long, repeated-prefix prompt twice to observe a cache hit on the second call (check router logs / metrics for prefix-cache reuse).

Troubleshooting

- **Pod fails to start under Pod Security Admission "restricted" policy**: the overlay requires `privileged: true`; either relax the namespace's PSA label or use a different offloading path that doesn't need privileged access (verify per-connector requirements before assuming).
- **No observable cache reuse**: confirm the CPU-tier offload size is large enough for your workload and that requests are actually re-sending the same prefix.

Cleanup

```
helm uninstall ${GUIDE_NAME} -n ${NAMESPACE}
kubectl delete -n ${NAMESPACE} -k
 ↪  ${REPO_ROOT}/guides/tiered-prefix-cache/modelserver/xpu/vllm/base
# or, if you deployed LMCache:
kubectl delete -n ${NAMESPACE} -k ${REPO_ROOT}/guides/tiered-prefix-
 ↪  cache/modelserver/xpu/vllm/lmcache-connector/cpu/base
```

Configuration Reference

Parameter	Value	Purpose
`deviceClassName`	`gpu.intel.com`	DRA device class (claim name `xpu-vllm-gpu`)
`securityContext.privileged`	`true`	Required for host-pinned memory management
CPU cache offload tier	CPU RAM (native) or CPU RAM (LMCache)	Effective KV-cache size beyond accelerator memory

---

1.5 Wide Expert Parallelism

Overview

Deploys a wide expert-parallel MoE model with P/D disaggregation using LeaderWorkerSets and DP-aware scheduling. The reference GPU configuration targets `DeepSeek-R1-0528` across 32 GPUs; the **validated Intel XPU configuration** uses the smaller `DeepSeek-V2-Lite-Chat` shape:

Parameter	Value
Model	`deepseek-ai/DeepSeek-V2-Lite-Chat`
Prefill Tensor Parallelism	2
Decode Tensor Parallelism	2
Total XPUs	4
Expert Parallelism	enabled
All2All backend	`allgather_reducescatter`
KV transfer	NIXL, `kv_buffer_device=xpu`
UCX transport	`tcp,ze_copy` (validated non-RDMA configuration)

> [!NOTE] The Intel XPU backend uses **XCCL** for collective communication; NIC-ID-based rail-only connectivity configurations (used on some NVIDIA RoCE setups) are not compatible and will fail.

Prerequisites

Chapter 0, plus the Intel GPU DRA driver already covers this case's device requirement — no additional driver install beyond Chapter 0.

Deployment Steps

Step 1 — Deploy the llm-d Router (Standalone mode, with the XPU override)

```
export GUIDE_NAME="wide-ep-lws"
export NAMESPACE="llm-d-wide-ep-lws"

helm install ${GUIDE_NAME} \
    ${ROUTER_STANDALONE_CHART} \
    -f ${REPO_ROOT}/guides/recipes/router/base.values.yaml \
    -f ${REPO_ROOT}/guides/${GUIDE_NAME}/router/${GUIDE_NAME}.values.yaml \
    -f ${REPO_ROOT}/guides/${GUIDE_NAME}/router/xpu.values.yaml \
    -n ${NAMESPACE} --version ${ROUTER_CHART_VERSION}
```

[!IMPORTANT] The `xpu.values.yaml` override is required — without it, EPP does not target the single decode sidecar port exposed by the Intel XPU manifests.

Step 2 — Deploy the Model Server

```
export MODEL=deepseek-ai/DeepSeek-V2-Lite-Chat
kubectl apply -n ${NAMESPACE} -k
 ↪  ${REPO_ROOT}/guides/wide-ep-lws/modelserver/xpu/vllm
```

Step 3 — (Optional) Enable Monitoring / Topology-Aware Scheduling

Monitoring follows the same pattern as other guides (`kubectl apply -k .../guides/recipes/modelserver/` pd, if applicable). TAS overlays in this guide are GKE H200/B200-specific and do not apply to the Intel XPU path.

Verification

```
kubectl get pods -n ${NAMESPACE}
kubectl get pods -n ${NAMESPACE} -l llm-d.ai/role=decode -w
kubectl get pods -n ${NAMESPACE} -l llm-d.ai/role=prefill -w
```

Inference Test

```
export IP=$(kubectl get service ${GUIDE_NAME}-epp -n ${NAMESPACE} -o
 ↪  jsonpath='{.spec.clusterIP}')

kubectl run curl-debug --rm -it \
    --image=cfmanteiga/alpine-bash-curl-jq \
    --namespace="${NAMESPACE}" \
    --env="IP=${IP}" \
    -- /bin/bash

curl -X POST http://${IP}/v1/completions \
    -H 'Content-Type: application/json' \
    -d '{"model": "deepseek-ai/DeepSeek-V2-Lite-Chat", "prompt": "Explain
     ↪  expert parallelism."}' | jq
```

Troubleshooting

- **Collective communication failures / hangs**: confirm the cluster is using XCCL, not a NIC-ID/rail-only RoCE configuration — the latter is documented as incompatible with the Intel XPU backend.

- **EPP routing to the wrong port**: confirm `router/xpu.values.yaml` was included in the `helm install` — without it EPP targets the wrong (non-XPU) sidecar port.

Cleanup

```
helm uninstall ${GUIDE_NAME} -n ${NAMESPACE}
kubectl delete -n ${NAMESPACE} -k
 ↪  ${REPO_ROOT}/guides/wide-ep-lws/modelserver/xpu/vllm
```

Configuration Reference

Parameter	Value	Purpose
`kv_buffer_device`	`xpu`	KV buffer resident on the accelerator (differs from the cpu default used by the PD-disaggregation guide)
All2All backend	`allgather_reducescatter`	MoE expert-parallel communication pattern
UCX transport	`tcp,ze_copy`	Validated non-RDMA transport for Intel XPU
Total XPUs	4	2 (prefill TP) + 2 (decode TP), validated shape

1.6 Multimodal Serving (Aggregated)

Overview

Deploys the recommended configuration for multimodal (image + text) vLLM serving with prefix-cache aware routing that matches combined text + image hashes, plus load-aware scoring. Default GPU model is `Qwen/Qwen3-VL-32B-Instruct`; the Intel XPU overlay targets **Intel Arc Pro B60**.

Prerequisites

Chapter 0 only.

Deployment Steps

Step 1 — Deploy the llm-d Router

```
export GUIDE_NAME="aggregation"
export NAMESPACE="llm-d-multimodal-aggregation"

helm install ${GUIDE_NAME} \
    ${ROUTER_STANDALONE_CHART} \
    -f ${REPO_ROOT}/guides/recipes/router/base.values.yaml \
    -f ${REPO_ROOT}/guides/multimodal-serving/${GUIDE_NAME}/router/
     ↪  ${GUIDE_NAME}.values.yaml \
    -n ${NAMESPACE} --version ${ROUTER_CHART_VERSION}
```

Step 2 — Deploy the Model Server (Intel XPU overlay)

```
kubectl apply -n ${NAMESPACE} -k
 ↪  ${REPO_ROOT}/guides/multimodal-serving/${GUIDE_NAME}/modelserver/xpu/vllm/
```

The overlay patches the decode Deployment's resource requests/limits and DRA claim (`intel-claim-template-decode`, requesting 1 `gpu.intel.com` device); the base image and multimodal serving args come from the shared base manifest.

Step 3 — (Optional) Enable Monitoring

```
kubectl apply -n ${NAMESPACE} -k
 ↪  ${REPO_ROOT}/guides/recipes/modelserver/components/monitoring
```

Verification

```
kubectl get pods -n ${NAMESPACE}
export IP=$(kubectl get service ${GUIDE_NAME}-epp -n ${NAMESPACE} -o
 ↪  jsonpath='{.spec.clusterIP}')
```

Inference Test

```
kubectl run curl-debug --rm -it \
    --image=cfmanteiga/alpine-bash-curl-jq \
    --namespace="${NAMESPACE}" \
    --env="IP=${IP}" \
    -- /bin/bash

curl -X POST http://${IP}/v1/chat/completions \
    -H 'Content-Type: application/json' \
    -d '{
      "model": "Qwen/Qwen3-VL-32B-Instruct",
      "messages": [{"role": "user", "content": [
        {"type": "text", "text": "What is in this image?"},
        {"type": "image_url", "image_url": {"url":
 ↪  "https://example.com/sample.jpg"}}
      ]}]
    }' | jq
```

Troubleshooting

- **GPU allocation fails**: confirm the node advertises `gpu.intel.com` resources compatible with Intel Arc Pro B60 and that the DRA driver version supports it.
- **Image decoding errors**: check vLLM logs for multimodal preprocessing errors; verify the image URL is reachable from within the cluster.

Cleanup

```
helm uninstall ${GUIDE_NAME} -n ${NAMESPACE}
kubectl delete -n ${NAMESPACE} -k
 ↪  ${REPO_ROOT}/guides/multimodal-serving/${GUIDE_NAME}/modelserver/xpu/vllm/
```

Configuration Reference

Parameter	Value	Purpose
`deviceClassName`	`gpu.intel.com`	DRA device class (claim `intel-claim-template-decode`)
Target hardware	Intel Arc Pro B60	As documented in the guide's hardware table

---

2.1 Flow Control

Overview

Hardware-agnostic — same steps apply on Intel XPU as any other accelerator.

Flow Control adds intelligent request queuing at the llm-d Router (EPP) level. Traditional load balancing falls short for LLM inference because resource consumption varies wildly per request; shifting queuing into the Router enables **multi-tenant fairness** (prevent noisy neighbors from starving other tenants) and **no-regret scheduling** (hold requests during saturation instead of committing them to a server's local queue where they get stuck). Source: `guides/flow-control/README.md`.

Requests are classified by a `FlowKey` (Fairness ID + Priority). The EPP maintains separate in-memory queues per flow and dispatches by priority band, then fairness across tenants within a band, then arrival order within a flow. Default policies (`global-strict-fairness-policy` + `fcfs-ordering-policy` + `utilization-detector`) intentionally mimic legacy FCFS behavior for a seamless transition — production deployments should switch the saturation detector to `concurrency-detector` per the guide's `tuning.md` to avoid telemetry lag risk.

Supported Hardware Backends: Flow Control is a software-level EPP feature and is entirely hardware-agnostic — it supports every accelerator the Optimized Baseline guide supports, including Intel XPU. This guide dynamically reuses whichever model server overlay you deploy from Optimized Baseline; there is no separate `modelserver/xpu/` overlay specific to Flow Control because none is needed.

Prerequisites

- Chapter 0 (Common Installation) complete.
- No additional cluster prerequisites beyond Chapter 0.

Deployment Steps

```
export REPO_ROOT=$(realpath $(git rev-parse --show-toplevel))
source ${REPO_ROOT}/guides/env.sh
export GUIDE_NAME="flow-control"
export NAMESPACE="llm-d-flow-control"
export MODEL_NAME="Qwen/Qwen3-32B"   # swap for the CI-sized Qwen/Qwen3-0.6B
 ↪  on constrained Intel XPU test nodes

kubectl create namespace ${NAMESPACE} --dry-run=client -o yaml | kubectl apply
 ↪  -f -

# HF token secret (see Chapter 0.4)
```

```
kubectl create secret generic llm-d-hf-token \
  --from-literal="HF_TOKEN=${HF_TOKEN}" \
  --namespace "${NAMESPACE}" --dry-run=client -o yaml | kubectl apply -f -

# 1. Deploy the Router (Standalone Mode)
helm install ${GUIDE_NAME} \
    ${ROUTER_STANDALONE_CHART} \
    -f ${REPO_ROOT}/guides/recipes/router/base.values.yaml \
    -f ${REPO_ROOT}/guides/${GUIDE_NAME}/router/${GUIDE_NAME}.values.yaml \
    -n ${NAMESPACE} --version ${ROUTER_CHART_VERSION}

# 2. Deploy the Model Server — reuse the Optimized Baseline Intel XPU overlay,
#    relabeled to this guide's name via `sed` (this is the pattern the guide
 ↪  itself
#    prescribes for every non-default accelerator, not an Intel-XPU-specific
 ↪  workaround)
kubectl kustomize ${REPO_ROOT}/guides/optimized-baseline/modelserver/xpu/vllm/
 ↪  \
  | sed "s/optimized-baseline/${GUIDE_NAME}/g" \
  | kubectl apply -n ${NAMESPACE} -f -
```

For Gateway Mode (Kubernetes Gateway rather than standalone Envoy sidecar), see `guides/flow-control/README.md` "Gateway Mode" — the pattern is identical to the one used in Chapter 1.

Verification

```
export IP=$(kubectl get service ${GUIDE_NAME}-epp -n ${NAMESPACE} -o
 ↪  jsonpath='{.spec.clusterIP}')
kubectl logs deploy/${GUIDE_NAME}-epp -n ${NAMESPACE} | grep "Flow Control
 ↪  enabled"
```

Inference Test

Apply the three priority-tier `InferenceObjectives` shipped with the guide, then send a request tagged with a tenant/priority header:

```
kubectl apply -f ${REPO_ROOT}/guides/${GUIDE_NAME}/objectives.yaml -n
 ↪  ${NAMESPACE}

kubectl run curl-debug --rm -it \
    --image=cfmanteiga/alpine-bash-curl-jq \
    --namespace="$NAMESPACE" \
    --env="IP=$IP" --env="NAMESPACE=$NAMESPACE" --env="GUIDE_NAME=$GUIDE_NAME"
     ↪  \
    -- /bin/bash

# from inside the debug pod:
curl -X POST http://${IP}/v1/completions \
  -H 'Content-Type: application/json' \
  -H 'x-llm-d-inference-fairness-id: tenant-a' \
  -H 'x-llm-d-inference-objective: premium-traffic' \
  -d "{\"model\": \"${MODEL_NAME}\", \"prompt\": \"Say hello\"}"

# confirm queuing metrics are exposed
```

```
curl http://${GUIDE_NAME}-epp:9090/metrics | grep
 ↪  llm_d_epp_flow_control_queue_size
```

To actually observe backpressure (not just admission), drive concurrent load with hey per the guide's "Use Case 2: Backpressure Management" section — a single request always dispatches immediately because the system is work-conserving.

Troubleshooting

- **Trust boundary**: never let end users self-assert `x-llm-d-inference-fairness-id` / `x-llm-d-inference-objective` directly — a production ingress Gateway must strip these headers from external traffic and re-inject them after authenticating the caller (see the EPP HTTP headers reference in `docs/api-reference/epp-http-headers.md`).
- If `"Flow Control enabled"` never appears in the EPP logs, confirm the `${GUIDE_NAME}.values.yaml` router overlay was actually applied (not just the base values file).
- `utilization-detector` (the default saturation detector) can lag under bursty Intel XPU telemetry collection intervals; switch to `concurrency-detector` per the Tuning Guide if you see saturation detected later than expected.

Cleanup

```
helm uninstall ${GUIDE_NAME} -n ${NAMESPACE}
# Must reuse the same kustomize+sed rename pipeline as the deploy step above —
 ↪  the resources were
# created under the ${GUIDE_NAME}-prefixed names, not the original
 ↪  optimized-baseline names, so a
# plain `kubectl delete -k .../optimized-baseline/modelserver/xpu/vllm/` would
 ↪  not match them.
kubectl kustomize ${REPO_ROOT}/guides/optimized-baseline/modelserver/xpu/vllm/
 ↪  \
  | sed "s/optimized-baseline/${GUIDE_NAME}/g" \
  | kubectl delete -n ${NAMESPACE} -f -
kubectl delete namespace ${NAMESPACE}
```

Config Reference

Field	Value	Notes
Default model (inherited)	`Qwen/Qwen3-32B`	8 replicas × TP2 = 16 GPUs in the reference config; scale down for Intel XPU test clusters
Fairness policy	`global-strict-fairness-policy`	ignores flow isolation, single global FCFS order
Ordering policy	`fcfs-ordering-policy`	first-come-first-served
Saturation detector	`utilization-detector` (default) / `concurrency-detector` (recommended for production)	
Priority tiers (example)	Premium (100), Standard (0), Best-Effort (-10)	defined in `objectives.yaml`

Field	Value	Notes
Client headers	`x-llm-d-inference-fairness-id`, `x-llm-d-inference-objective`	must be stripped/re-injected by ingress Gateway in production

2.2 Predicted Latency-Based Routing

Overview

Hardware-agnostic at the routing layer; the GPU/TPU-specific model-server paths shown upstream do not include an Intel XPU variant of their own — Intel XPU deployers reuse the Optimized Baseline overlay exactly as the guide's "For other backends" note instructs.

Routes each inference request to the model server predicted to serve it fastest, using a live-trained XGBoost latency-prediction model instead of heuristic queue-depth/KV-utilization scoring — and, optionally, only to a server predicted to meet a per-request TTFT/TPOT SLO. Source: `guides/predicted-latency-routing/README.md`.

When to pick this path: workloads with high variance in prompt/completion length where queue depth alone is a poor load proxy, or where clients can express per-request latency SLOs you want the gateway to enforce. **Skip it** for heterogeneous pools (mixed GPU types/model variants) — the predictor assumes a single pod shape and will produce inaccurate predictions otherwise.

> [!NOTE] Upstream flags OpenShift support for this guide as "not reliable as-is" — the latency-predictor sidecars may need additional OpenShift-specific runtime adjustments. Prefer a vanilla Kubernetes distribution for your first Intel XPU validation pass.

Prerequisites

- Chapter 0 (Common Installation) complete.
- No additional cluster prerequisites beyond Chapter 0.

Deployment Steps

```
export REPO_ROOT=$(realpath $(git rev-parse --show-toplevel))
source ${REPO_ROOT}/guides/env.sh
export GUIDE_NAME="predicted-latency-routing"
export NAMESPACE=llm-d-predicted-latency
export MODEL_NAME="Qwen/Qwen3-32B"

kubectl create namespace ${NAMESPACE}
kubectl create secret generic llm-d-hf-token \
  --from-literal="HF_TOKEN=${HF_TOKEN}" \
  --namespace "${NAMESPACE}" --dry-run=client -o yaml | kubectl apply -f -

# 1. Deploy the Router — default values file trains the predictor on
 ↪  end-to-end latency
#    (routing only, no SLO headers). Swap in
 ↪  router/predicted-latency-slo.values.yaml
#    for SLO-aware scheduling.
```

```
helm install ${GUIDE_NAME} \
    ${ROUTER_STANDALONE_CHART} \
    -f ${REPO_ROOT}/guides/recipes/router/base.values.yaml \
    -f ${REPO_ROOT}/guides/${GUIDE_NAME}/router/predicted-latency.values.yaml
     ↪  \
    -n ${NAMESPACE} --version ${ROUTER_CHART_VERSION}

# 2. Deploy the Model Server — Intel XPU is not one of this guide's own
 ↪  model-server
#    variants (only GPU/vLLM base+gke, and a TPU variant for the agentic case
 ↪  exist under
#    guides/predicted-latency-routing/modelserver/). Per the guide's own
 ↪  instructions for
#    "other backends", reuse the Optimized Baseline Intel XPU overlay directly
 ↪  — both
#    values files already select pods labeled
 ↪  `llm-d.ai/guide=optimized-baseline`, so no
#    relabeling `sed` step is required here (unlike Flow Control in Chapter
 ↪  2.1).
kubectl apply -n ${NAMESPACE} -k
 ↪  ${REPO_ROOT}/guides/optimized-baseline/modelserver/xpu/vllm/
```

Verification

Confirm predictions are actually being produced in Prometheus (see `docs/architecture/advanced/latency-predictor.md#observability` for the full metric reference):

`inference_objective_request_ttft_prediction_duration_seconds`

If this stays empty, the predictor sidecar isn't being called — tail the EPP logs for `predicted-latency-producer` errors.

Inference Test

```
export IP=$(kubectl get service ${GUIDE_NAME}-epp -n ${NAMESPACE} -o
 ↪  jsonpath='{.spec.clusterIP}')

kubectl run curl-debug --rm -it \
    --image=cfmanteiga/alpine-bash-curl-jq \
    --namespace="$NAMESPACE" \
    --env="IP=$IP" --env="NAMESPACE=$NAMESPACE" --env="MODEL_NAME=$MODEL_NAME"
     ↪  \
    -- /bin/bash

# from inside the debug pod — opt this request into SLO-aware routing:
curl -X POST http://${IP}/v1/completions \
    -H 'Content-Type: application/json' \
    -H 'x-llm-d-slo-ttft-ms: 200' \
    -H 'x-llm-d-slo-tpot-ms: 50' \
    -d '{
        "model": "'${MODEL_NAME}'",
        "prompt": "Explain the difference between prefill and decode.",
        "max_tokens": 200,
        "temperature": 0,
        "stream": true,
```

```
        "stream_options": {"include_usage": true}
    }'
```

No header changes are needed for plain latency-based routing — it applies to every request. SLO headers (`x-llm-d-slo-ttft-ms`, `x-llm-d-slo-tpot-ms`) opt a request into SLO enforcement; sheddable (negative-priority) requests are rejected at admission if no endpoint can meet the SLO, rather than routed to a guaranteed miss.

Troubleshooting

- **Empty prediction metrics**: predictor sidecar not being invoked — check EPP logs for `predicted-latency-producer` errors, and confirm the router values file actually enabled the predictor plugin (not just the base values file).
- **Predictions drifting from reality**: compare `inference_objective_request_predicted_ttft_secon` against `inference_objective_request_ttft_seconds` over a rolling window — a healthy deployment converges within a few percent after warmup; if it doesn't, the pool is likely heterogeneous (mixed hardware/model shapes), which this guide explicitly does not support.
- **OpenShift**: known-unreliable per upstream notes — try a vanilla Kubernetes distro first if you hit sidecar runtime issues.

Cleanup

```
helm uninstall ${GUIDE_NAME} -n ${NAMESPACE}
kubectl delete -n ${NAMESPACE} -k
 ↪  ${REPO_ROOT}/guides/optimized-baseline/modelserver/xpu/vllm/
kubectl delete namespace ${NAMESPACE}
```

Config Reference

Field	Value	Notes
Values file (default)	`router/predicted-latency.values.yaml`	routing-only, predictor trained on E2E latency
Values file (SLO-aware)	`router/predicted-latency-slo.values.yaml`	requires `"stream": true` on every request
SLO headers	`x-llm-d-slo-ttft-ms`, `x-llm-d-slo-tpot-ms`	opt-in per request
Model server (Intel XPU)	reuses `guides/optimized-baseline/modelserver/xpu/vllm/` directly	no relabeling needed, unlike Flow Control
Key metrics	`inference_objective_request_..._predicted_ttft_seconds`, `..._ttft_slo_violation_total`	see architecture doc for full list

2.3 Workload Autoscaling

Overview

Hardware-agnostic — scales whatever model-server Deployment you already run (including the Intel XPU Optimized Baseline overlay); no accelerator-specific autoscaling logic exists.

CPU/GPU utilization metrics are poor autoscaling signals for LLM inference — accelerators often sit near 100% "utilized" during active batching regardless of actual load. This guide covers two proactive, SLO-aware autoscaling paths built on real inference-demand signals (queue depth, in-flight requests, KV-cache pressure). Source: `guides/workload-autoscaling/README.md`.

Path	Signal source	Best for
KEDA + EPP Metrics (`README.hpa-epp.md`)	EPP-emitted queue depth / running-request-count metrics	Homogeneous hardware, each model-server pool scaled independently
HPA + WVA Metrics (`README.wva.md`)	Workload Variant Autoscaler's `wva_desired_replicas` (KV utilization, queue depth, perf budgets)	Multi-variant deployments across heterogeneous hardware, cost-aware scaling

An experimental **Replica Rebalancing** feature (`README.replica-rebalancing.md`) adjusts max replica counts across annotated HPAs sharing a GPU budget — explicitly flagged upstream as a proof-of-concept, not production-ready.

KEDA (or the OpenShift Custom Metrics Autoscaler Operator) is the recommended metrics adapter for both paths — Prometheus Adapter is deprecated. The **KEDA + EPP Metrics** path documents KEDA as the primary supported route. The **HPA + WVA Metrics** path can use KEDA-managed HPA (recommended) or a plain HPA reading the `wva_desired_replicas` external metric directly, without KEDA.

Prerequisites

- Chapter 0 (Common Installation) complete, plus a running **Prometheus** instance (see `docs/operations/observability/setup.md`) — WVA requires TLS enabled on Prometheus.
- **KEDA** installed (recommended path) — see the KEDA install guide. On OpenShift, use the Custom Metrics Autoscaler Operator instead.
- Complete Optimized Baseline on Intel XPU first, **including its optional monitoring step** — this guide layers on top of it and needs the EPP metrics endpoint already being scraped.

Deployment Steps (KEDA + EPP Metrics path)

```
export REPO_ROOT=$(realpath $(git rev-parse --show-toplevel))
source ${REPO_ROOT}/guides/env.sh
export NAMESPACE=llm-d-optimized-baseline
export MONITORING_NAMESPACE=llm-d-monitoring
export KEDA_NAMESPACE=keda

# Upgrade the Optimized Baseline router with the KEDA+EPP overlay (enables EPP
 ↪  Flow Control)
# — reapply the monitoring feature values so the EPP metrics
 ↪  port/ServiceMonitor stay enabled
helm upgrade optimized-baseline \
  ${ROUTER_STANDALONE_CHART} \
  -f ${REPO_ROOT}/guides/recipes/router/base.values.yaml \
  -f ${REPO_ROOT}/guides/optimized-baseline/router/optimized-
   ↪  baseline.values.yaml \
```

```
  -f ${REPO_ROOT}/guides/recipes/router/features/monitoring.values.yaml \
  -f ${REPO_ROOT}/guides/workload-autoscaling/keda-epp/router.values.yaml \
  -n ${NAMESPACE} --version ${ROUTER_CHART_VERSION}
```

The rest of the KEDA `ScaledObject` setup (targeting `llm_d_epp_flow_control_queue_size` and `llm_d_epp_request_running` via PromQL) is identical regardless of the accelerator backing the model-server Deployment — see `guides/workload-autoscaling/README.hpa-epp.md` for the full `ScaledObject` manifest and PromQL queries.

For the WVA path (multi-variant, cost-aware scaling across heterogeneous hardware), see `guides/workload-autoscaling/README.wva.md` — most relevant if you run Intel XPU alongside other accelerator types and want WVA to prefer the cheaper variant automatically.

Verification

```
kubectl logs deployment/optimized-baseline-epp -n ${NAMESPACE} | grep "Flow
 ↪  Control enabled"
kubectl get servicemonitor -n ${NAMESPACE}

# confirm EPP exposes the metrics directly
kubectl port-forward -n ${NAMESPACE} service/optimized-baseline-epp 9091:9090
 ↪  &
curl -s http://localhost:9091/metrics | grep -E
 ↪  'llm_d_epp_flow_control_queue_size|llm_d_epp_request_running'
```

Then, in your Prometheus query UI, confirm the metric is actually being scraped:

```
sum(llm_d_epp_flow_control_queue_size{namespace="llm-d-optimized-
 ↪  baseline",service="optimized-baseline-epp",model_name="Qwen/Qwen3-32B"})
```

Inference Test

Drive sustained concurrent load against the Optimized Baseline endpoint (see Optimized Baseline's Inference Test) and watch the `ScaledObject`'s managed HPA add replicas as queue depth/running-request-count rise, then scale back down as load subsides.

Troubleshooting

- **Do not create a separate HPA** for a Deployment already managed by a KEDA `ScaledObject` — two HPAs targeting the same Deployment produce conflicting scaling decisions. KEDA's generated HPA remains visible for inspection only.
- If metrics never populate in Prometheus, confirm the `ServiceMonitor` survived the router `helm upgrade` in the deployment step — reapplying the KEDA+EPP overlay without also re-including the monitoring values file will silently disable metric scraping.
- WVA requires TLS on the Prometheus endpoint; if the WVA controller can't reach Prometheus, check certificate/auth configuration first.

Cleanup

```
# Remove the KEDA ScaledObject (see README.hpa-epp.md for its exact name)
kubectl delete scaledobject <scaledobject-name> -n ${NAMESPACE}

# Revert the router to the plain Optimized Baseline values (drop the keda-epp
 ↪  overlay)
helm upgrade optimized-baseline \
```

```
${ROUTER_STANDALONE_CHART} \
-f ${REPO_ROOT}/guides/recipes/router/base.values.yaml \
-f ${REPO_ROOT}/guides/optimized-baseline/router/optimized-
 ↪  baseline.values.yaml \
-n ${NAMESPACE} --version ${ROUTER_CHART_VERSION}
```

Config Reference

Field	Value	Notes
Scaling signal (KEDA+EPP)	`llm_d_epp_flow_control_queue_size`, `llm_d_epp_request_running`	EPP native metrics
Scaling signal (WVA)	`wva_desired_replicas`	aggregates KV utilization + queue depth + perf budget
Scale-to-zero	Supported on both paths	
Additional components (KEDA path)	KEDA, Prometheus	no Prometheus Adapter needed
Additional components (WVA path)	WVA controller, Prometheus (TLS required)	
Replica Rebalancing	experimental, proof-of-concept only	not production-ready per upstream

2.4 Rollouts

> **Repo location note**: the guides-index (`guides/README.md`) links to `./rollouts/README.md`, but that path has moved to `docs/operations/rollouts/` as part of a broader documentation reorganization ("Reorganize part 2: create two new pillars, operations and infrastructure"). The commands below reference the current `docs/operations/rollouts/` path.

Overview

Hardware-agnostic — directly useful for migrating an existing NVIDIA deployment to Intel XPU (or vice versa) with minimal disruption, see "Node/Accelerator Update Roll Out" below.

Rollout guides cover three incremental deployment strategies for updating inference infrastructure with minimal service disruption. Source: `docs/operations/rollouts/README.md`.

Strategy	Mechanism	Deployment mode	Best for
Rolling Update	Standard Kubernetes Deployment rolling update (e.g. 25% at a time)	Standalone or Gateway	General, non-critical updates; conserves compute
Blue-Green Update	Second complete `InferencePool` + `HTTPRoute` traffic-weight splitting	Gateway only	Critical production rollouts, fast rollback, canary by header

Strategy	Mechanism	Deployment mode	Best for
LoRA Adapter Rollout	`InferenceModelRewrite` mapping model names to adapter versions	Standalone or Gateway	Updating LoRA adapters without touching base model/infra

Why this matters for Intel XPU specifically: Blue-Green Update's documented use case #1 is literally "**Node(compute, accelerator) update roll out**" — safely migrating inference workloads to new node hardware or accelerator configurations without interrupting service. This is the recommended mechanism for migrating an existing NVIDIA `InferencePool` to Intel XPU nodes incrementally (e.g. 1% → 10% → 50% → 100% traffic shift) rather than a hard cutover.

Prerequisites

- Chapter 0 (Common Installation) complete, with a working deployment already running (e.g. Optimized Baseline on Intel XPU) that you intend to update/migrate.
- For Blue-Green Update: llm-d Router in **Gateway mode** (a Kubernetes Gateway + HTTPRoute) — see `docs/infrastructure/gateway/`; Blue-Green does not work in standalone mode.

Deployment Steps

Rolling Update

Standard Kubernetes mechanics — update the vLLM image/tag or config in your existing Helm values/Kustomize overlay and re-apply; Kubernetes replaces pods incrementally while old pods keep serving traffic. No llm-d-specific steps beyond your normal `helm upgrade` / `kubectl apply -k`.

Blue-Green Update — example: Intel XPU node/accelerator migration

Starting from an existing `InferencePool` (e.g. `vllm-qwen3-32b` on NVIDIA GPU nodes):

```
# 1. Deploy a second, complete InferencePool (the "green" pool) on Intel XPU
 ↪  nodes,
#    with a different Helm release name / selector — e.g. by applying the
 ↪  Optimized
#    Baseline Intel XPU overlay under a new name:
kubectl kustomize ${REPO_ROOT}/guides/optimized-baseline/modelserver/xpu/vllm/
 ↪  \
  | sed "s/optimized-baseline/vllm-qwen3-32b-new/g" \
  | kubectl apply -n ${NAMESPACE} -f -

# 2. Edit the HTTPRoute to split traffic between old (blue, NVIDIA) and new
 ↪  (green, Intel XPU)
kubectl edit httproute llm-route
```

```
apiVersion: gateway.networking.k8s.io/v1
kind: HTTPRoute
metadata:
  name: llm-route
spec:
  parentRefs:
```

```
    - group: gateway.networking.k8s.io
      kind: Gateway
      name: inference-gateway
  rules:
    - backendRefs:
        - group: inference.networking.k8s.io
          kind: InferencePool
          name: vllm-qwen3-32b       # old (blue)
          weight: 90
        - group: inference.networking.k8s.io
          kind: InferencePool
          name: vllm-qwen3-32b-new   # new (green, Intel XPU)
          weight: 10
      matches:
        - path:
            type: PathPrefix
            value: /
```

Gradually increase the green weight (10 → 50 → 100) as you validate correctness/performance on Intel XPU, then finish the rollout by setting the green pool's weight to 100 and removing the `backendRefs` entry for the old pool entirely. Roll back instantly at any point by flipping the weights back.

LoRA Adapter Rollout

See `docs/operations/rollouts/adapter-rollout.md` — uses `InferenceModelRewrite` to gradually shift traffic between adapter versions without any infrastructure change. Fully accelerator agnostic; identical on Intel XPU.

Verification

```
IP=$(kubectl get gateway/inference-gateway -o
 ↪  jsonpath='{.status.addresses[0].value}'); PORT=80

# send a batch of requests and confirm the observed split roughly matches the
 ↪  configured weights
for i in $(seq 1 20); do
  curl -s -o /dev/null -w "%{http_code}\n" ${IP}:${PORT}/v1/completions \
    -H 'Content-Type: application/json' \
    -d '{"model":"vllm-qwen3-32b","prompt":"ping","max_tokens":1}'
done
```

Cross-check which pool actually served each request via response headers or per-pool EPP logs.

Inference Test

```
curl -i ${IP}:${PORT}/v1/completions -H 'Content-Type: application/json' -d '{
  "model": "vllm-qwen3-32b",
  "prompt": "Write as if you were a critic: San Francisco",
  "max_tokens": 100,
  "temperature": 0
}'
```

Troubleshooting

- Blue-Green **requires Gateway mode** — attempting this in standalone mode has no HTTPRoute to edit; use Rolling Update or LoRA Adapter Rollout instead if you're on standalone.
- Keep the old (blue) `InferencePool` and its nodes running until the new (green) pool is fully validated — this is your rollback path; deleting it early forfeits instant rollback.
- If traffic doesn't appear to split according to configured weights, confirm both `InferencePools` are correctly referenced in `backendRefs` and that neither pool's own Helm chart was installed with `httpRoute.create=true` (a pool-level catch-all route will conflict with this guide's manually-edited HTTPRoute, the same conflict called out in Multi-Model Routing).

Cleanup

```
# Once fully migrated to the new (green) pool, remove the old one:
helm uninstall vllm-qwen3-32b -n ${NAMESPACE}
```

Config Reference

Field	Value	Notes
Strategy	Blue-Green (HTTPRoute weight-based traffic split)	Gateway mode only
`backendRefs[].weight`	integer, relative	e.g. 90/10 → 50/50 → 0/100
Rollback	flip weights back to 100/0	instant, no pod recreation needed
Alternative strategy	Rolling Update	standalone + gateway, in-place pod replacement
Alternative strategy	LoRA Adapter Rollout	`InferenceModelRewrite`, no infra change

3.1 Agentic Serving

Overview

This is a **workload composition guide**, not a standalone deployable overlay — its two ready-made hardware-specific deployments (NVIDIA H200, Google TPU v7) do not include an Intel XPU variant; Intel XPU users compose the same building blocks manually as described below.

Agentic Serving is a horizontal, workload-centric umbrella guide for serving agentic *programs* (e.g. coding agents) on llm-d — it composes several capability guides into one recommended, cohesive deployment rather than being a single new feature. Source: `guides/agentic-serving/README.md`.

The reference workload is **long-horizon loops** (agentic code generation): deep multi-turn sessions over large, repository-scale contexts with tool-call pauses between turns. Three behaviors drive the design: prefill-heavy/decode-light (large context dominates TTFT), high reusable locality (cache hit rate — not FLOPs — sets throughput), and bursty/stateful arrivals (tool pauses leave sessions idle, then resume in bursts).

The optimization stack (each layer already covered elsewhere in this whitepaper):

Layer	Chapter	What it does for the workload
Optimized Baseline	1.1	Prefix-cache scorer routes a turn to the replica already holding its prefix
Tiered Prefix Cache	1.4	Offloads KV cache beyond accelerator memory so idle sessions restore on resume instead of recomputing prefill
Precise Prefix Cache Routing	1.3	Exact, global cache-state view enabling session-centric orchestration
P/D Disaggregation	1.2	Separates prefill/decode pools so heavy prefill never stalls token generation

Upstream ships two **ready-made** deployments composing this stack, neither of which targets Intel XPU:

- `nemotron-3-ultra-550b-h200.md` — P/D-disaggregated serving on 8× NVIDIA H200 with CPU KV-offloading and ready-to-use coding-agent client configs.
- `qwen3-coder-480b-tpu.md` — routing + CPU KV-offloading on 8× Google TPU v7x (2x2x1).

Prerequisites

- Chapter 0 (Common Installation) complete.
- Each of the four composed guides' own prerequisites (Chapters 1.1–1.4 of this whitepaper), since Agentic Serving is additive on top of them rather than a replacement.

Deployment Steps (Intel XPU — compose the building blocks manually)

Since no ready-made Intel XPU deployment ships upstream, assemble the stack yourself from the already-validated Chapter 1 building blocks, following the same pattern as the two upstream examples (routing + KV-offloading, optionally + P/D disaggregation for larger models):

```
# 1. Deploy P/D Disaggregation on Intel XPU as the serving foundation for a
 ↪  large model
#    (see Chapter 1.2 for full steps) — this gives you separate prefill/decode
 ↪  pools.

# 2. Layer Tiered Prefix Cache's KV-offloading connector on top of the same
 ↪  model server
#    (see Chapter 1.4) so idle agentic sessions can resume from offloaded KV
 ↪  state instead of
#    recomputing the full prompt prefix.

# 3. Deploy the router with BOTH the P/D and tiered-prefix-cache values files
 ↪  layered together
#    (adjust chart/values paths to match your actual guide combination):
helm install agentic-xpu \
    ${ROUTER_STANDALONE_CHART} \
    -f ${REPO_ROOT}/guides/recipes/router/base.values.yaml \
    -f ${REPO_ROOT}/guides/pd-disaggregation/router/pd-
     ↪  disaggregation.values.yaml \
```

```
-f ${REPO_ROOT}/guides/tiered-prefix-cache/router/tiered-prefix-
 ↪  cache.values.yaml \
-n ${NAMESPACE} --version ${ROUTER_CHART_VERSION}
```

Caveat: this composition has not been validated by upstream CI for Intel XPU (unlike the individual Chapter 1 guides, which do have Intel XPU CI or documented validated configs). Treat this as a starting point requiring your own end-to-end validation, not a tested recipe.

Verification

Reuse each composed guide's own verification steps (P/D disaggregation pod health, tiered prefix-cache offload connector health) — see Chapters 1.2 and 1.4.

Inference Test

Use a realistic agentic-style prompt sequence (large repository-scale context, multiple turns reusing the same prefix) rather than a single short completion, since the whole point of this composition is prefix-cache reuse across turns:

```
curl -X POST http://${IP}/v1/chat/completions \
  -H 'Content-Type: application/json' \
  -d '{"model": "'${MODEL_NAME}'", "messages": [
        {"role": "system", "content": "<large repository context>"},
        {"role": "user", "content": "Refactor this function to be async"}
      ], "max_tokens": 500}'
```

Troubleshooting

- If cache-reuse benefits don't materialize, confirm Precise Prefix Cache Routing's `--block-size` matches the router's `tokenProcessorConfig.blockSize` exactly (both default 64 — see Chapter 1.3) — a mismatch silently degrades cache-aware scoring rather than erroring out.
- If offloaded sessions don't resume correctly, check that Tiered Prefix Cache's `securityContext.privileged: true` requirement (Chapter 1.4) is actually satisfied under your cluster's Pod Security Admission policy.

Cleanup

```
helm uninstall agentic-xpu -n ${NAMESPACE}
# then clean up each composed guide's model-server overlay per its own Chapter
 ↪  1 cleanup steps
```

Config Reference

Field	Value	Notes
Composed layers	Optimized Baseline + Tiered Prefix Cache + Precise Prefix Cache Routing + P/D Disaggregation	

Field	Value	Notes
Ready-made deployments (upstream)	NVIDIA H200 (Nemotron-3-Ultra-550B), Google TPU v7 (Qwen3-Coder-480B)	neither targets Intel XPU
Intel XPU status	Manual composition of Chapter 1 building blocks	not upstream-CI-validated as a combined stack
Benchmark harness	`inference-perf` via `llm-d-benchmark`, replaying agentic-style traces	see guide for exact preset

3.2 Multi-Model Routing

Overview

Hardware-agnostic — the Inference Payload Processor (IPP) and HTTPRoute-based routing operate purely at the Gateway layer; each backing `InferencePool` can independently be an Intel XPU deployment.

Deploys the **Inference Payload Processor (IPP)** to serve multiple LLMs behind a single Gateway endpoint. IPP extracts the model name from the request body and sets routing headers; HTTPRoutes then match those headers to direct traffic to the correct `InferencePool`. Source: `guides/multi-model-routing/README.md`.

Use this guide when you need to serve multiple base models (e.g. Qwen for chat, DeepSeek for reasoning) behind one OpenAI-compatible endpoint. For a single model, use Optimized Baseline instead.

Prerequisites

- Chapter 0 (Common Installation) complete.
- **Multiple `InferencePools` already deployed**, each serving a different base model — follow Optimized Baseline independently for each pool you want to add (e.g. once on Intel XPU for Qwen, once for a second model). When deploying each pool for this guide, do **not** pass `--set httpRoute.create=true` — this guide's own HTTPRoutes (deployment step 3) handle model-name-header-based routing, and a pool-level catch-all route would conflict.
- A Kubernetes Gateway (Istio, GKE, AgentGateway, etc.) — see `docs/infrastructure/gateway/README.`

Deployment Steps

```
export REPO_ROOT=$(realpath $(git rev-parse --show-toplevel))
source ${REPO_ROOT}/guides/env.sh
export GUIDE_NAME="multi-model-routing"
export NAMESPACE="llm-d-multi-model"

kubectl create namespace ${NAMESPACE} --dry-run=client -o yaml | kubectl apply
 ↪  -f -

# 1. Deploy IPP (clone its own repo; it is not part of the llm-d repo)
git clone https://github.com/llm-d/llm-d-inference-payload-processor.git
 ↪  /tmp/ipp
```

```
helm install ipp /tmp/ipp/config/charts/payload-processor \
    --set provider.name=istio \
    --set inferenceGateway.name=llm-d-inference-gateway \
    --set payloadProcessor.image.tag=v0.1.0-rc.4 \
    -n ${NAMESPACE}
# --set provider.name=gke on GKE. --set provider.name=none skips the
 ↪  auto-generated
# Istio EnvoyFilter/GKE GCPRoutingExtension — a Gateway is still required
 ↪  either way, you'd just
# wire IPP's ext-proc integration into it yourself instead of letting the
 ↪  chart do it.

kubectl get pods -n ${NAMESPACE} -l app=payload-processor

# 2. Register each base model with IPP via ConfigMap (edit
 ↪  manifests/configmaps.yaml to match
#    your actual Intel XPU model names/pools first)
kubectl apply -n ${NAMESPACE} -f
 ↪  ${REPO_ROOT}/guides/multi-model-routing/manifests/configmaps.yaml
# each ConfigMap needs label inference.llm-d.ai/ipp-managed: "true" and a
 ↪  baseModel matching
# your Intel XPU pool's actual served model name

# 3. Configure HTTPRoutes matching IPP's injected X-Gateway-Base-Model-Name
 ↪  header
#    (edit manifests/httproutes.yaml: parentRefs -> your Gateway, backendRefs
 ↪  -> your pool names)
kubectl apply -n ${NAMESPACE} -f
 ↪  ${REPO_ROOT}/guides/multi-model-routing/manifests/httproutes.yaml
```

Verification

```
kubectl logs -n ${NAMESPACE} -l app=payload-processor --tail=100
kubectl get configmap -l inference.llm-d.ai/ipp-managed=true -n ${NAMESPACE}
# confirm requests reach the correct pool (replace <pool-name> with your
 ↪  InferencePool name)
kubectl logs -n ${NAMESPACE} -l llm-d-router-gateway=<pool-name>-epp --tail=20
```

Inference Test

```
export GATEWAY_IP=$(kubectl get gateway llm-d-inference-gateway -n
 ↪  ${NAMESPACE} -o jsonpath='{.status.addresses[0].value}')

# Request routed to your first (e.g. Intel XPU Qwen) pool
curl -X POST "http://${GATEWAY_IP}/v1/chat/completions" \
  -H "Content-Type: application/json" \
  -d '{"model": "Qwen/Qwen3-32B", "messages": [{"role": "user", "content":
   ↪  "Hello"}]}'

# Request routed to a second pool
curl -X POST "http://${GATEWAY_IP}/v1/chat/completions" \
  -H "Content-Type: application/json" \
  -d '{"model": "deepseek/DeepSeek-r1", "messages": [{"role": "user",
   ↪  "content": "Solve this problem"}]}'
```

Advanced: LoRA Adapter Routing

Once base model routing works, extend a ConfigMap with an `adapters` field to route LoRA adapter names to their base model's pool:

```
kubectl apply -n ${NAMESPACE} -f - <<EOF
apiVersion: v1
kind: ConfigMap
metadata:
  name: qwen-model-mapping
  labels:
    inference.llm-d.ai/ipp-managed: "true"
data:
  baseModel: "Qwen/Qwen3-32B"
  adapters: |
    - food-review-1
    - travel-assistant
EOF

curl -X POST "http://${GATEWAY_IP}/v1/chat/completions" \
  -H "Content-Type: application/json" \
  -d '{"model": "food-review-1", "messages": [{"role": "user", "content":
   ↪  "Review this restaurant"}]}'
```

All model and adapter names must be globally unique across all `InferencePools`/ConfigMaps.

Troubleshooting

- Requests routed to the wrong pool (or 404): confirm the ConfigMap's `baseModel` value matches the model name clients send exactly, and that the `inference.llm-d.ai/ipp-managed: "true"` label is present — IPP logs discovered model mappings at startup, check them first.
- If a pool's traffic bypasses IPP header-based routing entirely, verify that pool's Helm install did **not** include `--set httpRoute.create=true` (see Prerequisites) — its own catch-all route would out-compete this guide's header-matched HTTPRoutes.

Cleanup

```
kubectl delete -n ${NAMESPACE} -f
 ↪  ${REPO_ROOT}/guides/multi-model-routing/manifests/httproutes.yaml
kubectl delete -n ${NAMESPACE} -f
 ↪  ${REPO_ROOT}/guides/multi-model-routing/manifests/configmaps.yaml
helm uninstall ipp -n ${NAMESPACE}
rm -rf /tmp/ipp
kubectl delete namespace ${NAMESPACE}
```

Config Reference

Field	Value	Notes
Routing header	`X-Gateway-Base-Model-Name`	injected by IPP, matched by HTTPRoute
ConfigMap label	`inference.llm-d.ai/ipp-managed: "true"`	required for IPP discovery

Field	Value	Notes
Per-pool Helm flag	must NOT set `httpRoute.create=true`	avoids catch-all route conflicts
LoRA routing	`adapters` field in ConfigMap	maps adapter name -> base model's pool
IPP repo	`github.com/llm-d/llm-d-inference-payload-processor`	separate repo, not part of `llmd`

4.1 Asynchronous Processing

Overview

Hardware-agnostic — the Async Processor dispatches to whatever llm-d Router endpoint you give it; the backing model server can be any Optimized Baseline overlay, including Intel XPU.

The Async Processor processes inference requests asynchronously via a queue-based architecture — ideal for latency-insensitive workloads or filling idle accelerator capacity. It decouples submission from execution (clients submit to a queue, retrieve results later), optimizes resource utilization, and provides automatic retries without impacting real-time traffic. Source: `guides/asynchronous-processing/README.md`.

Two supported queue implementations:

- **GCP Pub/Sub** (`gcp-pubsub/README.md`) — cloud-native, scalable messaging.
- **Redis Sorted Set** (`redis/README.md`) — high-performance, persisted, prioritized queue.

Prerequisites

- Kubernetes v1.31+ (Kind/Minikube for local dev; GKE/AKS/OpenShift for production).
- **Gateway Mode only**: a Gateway control plane deployed (`docs/infrastructure/gateway/README.md`). Not needed if you use Standalone Mode (step 1 below), which talks to the Router's Service `clusterIP` directly instead of a Gateway.
- An existing Optimized Baseline stack on Intel XPU to dispatch requests to — Async Processor is additive, not a replacement for the model-serving stack.

Deployment Steps

```
export REPO_ROOT=$(realpath $(git rev-parse --show-toplevel))

# 1. Get the IP of the llm-d Router deployed for Optimized Baseline (Chapter
 ↪  1.1)
# Standalone Mode:
export IP=$(kubectl get service optimized-baseline-epp -n
 ↪  llm-d-optimized-baseline -o jsonpath='{.spec.clusterIP}')
# Gateway Mode:
export IP=$(kubectl get gateway llm-d-inference-gateway -n
 ↪  llm-d-optimized-baseline -o jsonpath='{.status.addresses[0].value}')

# 2. Choose a queue implementation and review/edit its values.yaml:
#    guides/asynchronous-processing/gcp-pubsub/values.yaml
#    guides/asynchronous-processing/redis/values.yaml
```

```
# 3. Deploy the Async Processor
export NAMESPACE=llm-d-async
export MQ_PROVIDER=redis   # or gcp-pubsub
export ASYNC_VERSION=0.6.1

helm install async-processor \
    oci://ghcr.io/llm-d-incubation/charts/async-processor \
    -f ${REPO_ROOT}/guides/asynchronous-processing/${MQ_PROVIDER}/values.yaml
     ↪  \
    --set ap.igwBaseURL=http://${IP}:80 \
    -n ${NAMESPACE} --create-namespace --version ${ASYNC_VERSION}
```

Verification

```
kubectl get pods -n ${NAMESPACE}
kubectl logs -n ${NAMESPACE} -l app.kubernetes.io/name=async-processor
 ↪  --tail=50
```

Inference Test

Testing steps are queue-implementation specific — follow the "Testing" section of whichever sub-guide you chose:

- Redis: `guides/asynchronous-processing/redis/README.md#testing`
- GCP Pub/Sub: `guides/asynchronous-processing/gcp-pubsub/README.md#testing`

Both follow the same submit-then-poll pattern: publish a request payload to the configured queue, then poll for the corresponding result once the Async Processor has dispatched it to the Intel XPU-backed Optimized Baseline endpoint and received a response.

Troubleshooting

- If the Async Processor never dispatches requests, confirm `ap.igwBaseURL` actually resolves to a live llm-d Router endpoint (`curl` it directly first) — a common mistake is pointing at the wrong namespace's Router IP.
- Redis is documented as sufficient for both development and production priority-queue needs; GCP Pub/Sub requires GCP-side IAM/topic setup not covered by the llm-d chart itself.

Cleanup

```
helm uninstall async-processor -n ${NAMESPACE}
```

Config Reference

Field	Value	Notes
Queue implementations	GCP Pub/Sub, Redis Sorted Set	choose one via `MQ_PROVIDER`
Dispatch target	`ap.igwBaseURL`	must point at a live llm-d Router (any accelerator, incl. Intel XPU)

Field	Value	Notes
Chart	`oci://ghcr.io/llm-d-incubation/charts/async-processor`	separate incubation repo

4.2 Batch Gateway

Overview

Hardware-agnostic — Batch Gateway dispatches to an existing llm-d Router endpoint; the backing model server can be any Optimized Baseline overlay, including Intel XPU.

Batch Gateway provides an OpenAI-compatible Batch API (`/v1/batches`, `/v1/files`) for submitting, tracking, and managing large-scale batch inference jobs (up to 50,000 requests per job, configurable), designed to coexist with interactive workloads on shared infrastructure. Source: `guides/batch-gateway/README.md`.

Three components: an **API Server** (REST endpoints), a **Batch Processor** (pulls jobs from a priority queue, builds per-model execution plans, dispatches to the llm-d Router, writes output files), and a **Garbage Collector** (cleans up expired jobs/files).

Prerequisites

- Kubernetes v1.25+, Helm 3.0+.
- An existing Optimized Baseline stack on Intel XPU to dispatch requests to.
- **PostgreSQL 12+** for jobs/files metadata (Redis/Valkey dev-only alternative).
- **Redis 6+ / Valkey 8+** for the priority queue, events, and status updates.
- **S3 or a Filesystem PVC** for batch input/output file storage.

Deployment Steps

```
export NAMESPACE=batch-gateway
kubectl create namespace ${NAMESPACE}

# 1. Storage/queue credentials secret
kubectl create secret generic batch-gateway-secrets -n ${NAMESPACE} \
  --from-literal=redis-url="redis://redis-
   ↪  master.redis.svc.cluster.local:6379/0" \
  --from-literal=postgresql-url=
   ↪  "postgresql://user:pass@postgresql.postgresql.svc.cluster.local:5432/batchgateway"
   ↪  \
  --from-literal=s3-secret-access-key="<your-s3-secret-key>"

# 2. Point the Batch Processor at your Intel XPU-backed llm-d Router
export INFERENCE_GW_URL="http://infra-inference-scheduling-inference-gateway-
 ↪  istio.llm-d-inference-scheduler.svc.cluster.local:80"

# 3. Deploy
helm install batch-gateway oci://ghcr.io/llm-d-incubation/charts/batch-gateway
 ↪  \
  -n ${NAMESPACE} \
  --set processor.config.globalInferenceGateway.url="${INFERENCE_GW_URL}" \
```

```
  --set "apiserver.config.batchAPI.passThroughHeaders={Authorization}" \
  --set global.fileClient.fs.pvcName="batch-gateway-pvc"
```

For per-model gateways (different Intel XPU pools per model) instead of one global gateway URL, use `processor.config.modelGateways` — see the Helm chart README.

Verification

```
kubectl get pods -n ${NAMESPACE}   # expect apiserver, processor, gc all
 ↪  Running
kubectl port-forward -n ${NAMESPACE} svc/batch-gateway-apiserver 8081:8081 &
curl http://localhost:8081/health
```

`batch-gateway-apiserver` exposes two separate ports on the same Service: `8081` for the health/readiness endpoint above, and `8000` for the actual REST API (confirmed by the in-cluster URL `http://batch-gateway-apiserver:8000` used in the project's own benchmark manifests). The Inference Test below needs its own port-forward to `8000`.

Inference Test

```
kubectl port-forward -n ${NAMESPACE} svc/batch-gateway-apiserver 8000:8000 &

# 1. Prepare a JSONL input file (OpenAI Batch API format)
cat > batch_input.jsonl <<'EOF'
{"custom_id": "req-001", "method": "POST", "url": "/v1/chat/completions",
 ↪  "body": {"model": "my-model", "messages": [{"role": "user", "content":
 ↪  "Hello"}], "max_tokens": 100}}
{"custom_id": "req-002", "method": "POST", "url": "/v1/chat/completions",
 ↪  "body": {"model": "my-model", "messages": [{"role": "user", "content":
 ↪  "What is llm-d?"}], "max_tokens": 200}}
EOF

# 2. Upload it
curl -X POST http://localhost:8000/v1/files -F "purpose=batch" -F
 ↪  "file=@batch_input.jsonl"

# 3. Create the batch job
curl -X POST http://localhost:8000/v1/batches \
  -H "Content-Type: application/json" \
  -d '{"input_file_id": "<file-id-from-upload>", "endpoint":
   ↪  "/v1/chat/completions", "completion_window": "24h"}'

# 4. Monitor
curl http://localhost:8000/v1/batches/<batch-id> | jq '{status,
 ↪  request_counts}'

# 5. Download results once status is "completed"
OUTPUT_FILE_ID=$(curl -s http://localhost:8000/v1/batches/<batch-id> | jq -r
 ↪  '.output_file_id')
curl http://localhost:8000/v1/files/${OUTPUT_FILE_ID}/content > results.jsonl
```

Troubleshooting

- `passThroughHeaders` must include any auth header (e.g. `Authorization`) your llm-d Router expects — the Batch Processor forwards these when dispatching individual re-

quests; omitting a required header causes silent per-request auth failures inside batch jobs.
- For a production deployment with authentication, authorization, and TLS, see the Kubernetes deployment guide (Istio + Kuadrant + cert-manager) rather than the quick-start Helm install above.

Cleanup

```
helm uninstall batch-gateway -n ${NAMESPACE}
kubectl delete namespace ${NAMESPACE}
```

Config Reference

Field	Value	Notes
Max requests/job	50,000 (configurable)	
Metadata storage	PostgreSQL (prod) / Redis-Valkey (dev/test only)	
Queue/events storage	Redis/Valkey	
File storage	S3 or Filesystem PVC	
Dispatch target	`processor.config.globalInferenceGateway.url` or per-model `modelGateways`	any accelerator, incl. Intel XPU
Related guide	Asynchronous Processing	per-request async queue; complementary to Batch Gateway's job-oriented API

5.1 No-Kubernetes Deployment

Overview

This guide targets NVIDIA GPUs + vLLM specifically — the model-server image, CLI flags, and `--gpus` container invocation shown upstream are NVIDIA-specific. It is **not** a maintained, first-class Intel XPU path today; Intel XPU support is a manual substitution (detailed below), not something upstream provides or tests.

Deploys the llm-d routing stack (EPP + Envoy + one or more vLLM workers) **without a Kubernetes cluster** — the EPP gets its endpoint inventory from a plain YAML file on disk via the file-discovery plugin instead of watching a Kubernetes `InferencePool`. Configs are plain YAML — no Helm chart, no Kustomize overlay; drop them on a host and run. Source: `guides/no-kubernetes-deployment/README.md`.

The EPP and Envoy configs are themselves **accelerator-agnostic** — only the model-server container invocation (image, `--gpus` flag) is NVIDIA-specific.

What Needs to Change for Intel XPU

Component	Upstream (NVIDIA)	Intel XPU substitution
Model server image	`vllm/vllm-openai:v0.19.1`	An Intel XPU-built vLLM image (`VLLM_TARGET_DEVICE=xpu`) — see the image-build notes referenced from Optimized Baseline / the equivalent Intel XPU vLLM container build process used for the Kustomize overlays in Chapter 1
GPU allocation flag	`docker run --gpus '"device=0,1"'`	Intel XPU has no direct Docker `--gpus` equivalent; instead expose the render/video device nodes: `docker run --device=/dev/dri:/dev/dri` and add the container user to the host's `render`/`video` groups (same GIDs used in the Kustomize overlays' commented-out `securityContext.supplementalGroups`, typically 44/991 — verify on your actual host with `getent group render video`)
Tensor-parallel comms backend	NCCL (`--shm-size=20g` for `/dev/shm`)	XCCL for Intel XPU multi-device collectives (see the Wide-EP guide for the XCCL note) — keep `--shm-size` sized similarly since vLLM's shared-memory IPC mechanism is backend-agnostic
EPP / Envoy config	unchanged	fully accelerator-agnostic, no changes needed

Prerequisites

- A host with Intel XPU device(s) exposed via `/dev/dri`, Docker (or Podman), and an Intel XPU vLLM image already built (see substitution table above).
- A HuggingFace token in `HUGGING_FACE_HUB_TOKEN`.
- Same environment setup as upstream:

```
export REPO_ROOT=$(realpath $(git rev-parse --show-toplevel))
source ${REPO_ROOT}/guides/env.sh
export ENVOY_IMAGE=docker.io/envoyproxy/envoy:distroless-v1.33.2
export VLLM_IMAGE=<your-intel-xpu-vllm-image>   # substitution — not the
↪  upstream default
export MODEL=Qwen/Qwen3-32B
```

Deployment Steps

```
# 1. Stage the configs
sudo mkdir -p /etc/epp /etc/envoy
sudo cp guides/no-kubernetes-deployment/router/epp/config.yaml
 ↪  /etc/epp/config.yaml
sudo cp guides/no-kubernetes-deployment/router/epp/endpoints.yaml
 ↪  /etc/epp/endpoints.yaml
sudo cp guides/no-kubernetes-deployment/router/envoy/envoy.yaml
 ↪  /etc/envoy/envoy.yaml

# 2. Start the model server — Intel XPU substitution of the upstream `docker
 ↪  run --gpus` example
docker run -d --name vllm-0 \
    --device=/dev/dri:/dev/dri \
    --group-add "$(getent group render | cut -d: -f3)" \
    --shm-size=20g \
    -p 8000:8000 \
    -e "HUGGING_FACE_HUB_TOKEN=${HUGGING_FACE_HUB_TOKEN}" \
    -e "VLLM_TARGET_DEVICE=xpu" \
    -v "${HOME}/.cache/huggingface:/root/.cache/huggingface" \
    --entrypoint vllm \
    "${VLLM_IMAGE}" \
    serve "${MODEL}" \
    --disable-access-log-for-endpoints=/health,/metrics,/v1/models \
    --tensor-parallel-size=2

until curl -sf http://127.0.0.1:8000/v1/models >/dev/null; do sleep 2; done

# 3. Edit /etc/epp/endpoints.yaml to list each worker (literal IPv4 addresses
 ↪  required —
#    the file-discovery plugin does not resolve hostnames)

# 4. Start the EPP (accelerator-agnostic — unchanged from upstream)
docker run -d --name epp --network host \
    -v /etc/epp:/etc/epp:ro \
    "${ROUTER_EPP_IMAGE}:${ROUTER_EPP_VERSION}" \
    --config-file=/etc/epp/config.yaml \
    --pool-name=file-discovery --pool-namespace=default \
    --grpc-port=9002 --grpc-health-port=9003 --metrics-port=9090 \
    --secure-serving=false --v=2

# 5. Start Envoy (accelerator-agnostic — unchanged from upstream)
docker run -d --name envoy --network host \
    -v /etc/envoy/envoy.yaml:/etc/envoy/envoy.yaml:ro \
    "${ENVOY_IMAGE}" \
    --service-node envoy-proxy --log-level warn --concurrency 8 \
    --drain-strategy immediate --drain-time-s 60 \
    -c /etc/envoy/envoy.yaml
```

Verification

```
curl -s http://127.0.0.1:19000/ready
curl -s http://127.0.0.1:19000/clusters | grep -E
 ↪  '^(ext_proc|original_destination_cluster)'
```

```
curl -s http://127.0.0.1:9090/metrics | head
```

Inference Test

```
curl -s http://127.0.0.1:8081/v1/completions \
    -H 'Content-Type: application/json' \
    -d '{"model": "Qwen/Qwen3-32B", "prompt": "How are you today?"}'
```

Troubleshooting

- **EPP not detecting endpoint changes**: confirm `watchFile: true` in `router/epp/config.yaml`; update `endpoints.yaml` via atomic rename (`mv endpoints.yaml.tmp endpoints.yaml`), and check `docker logs epp` for `endpoints file changed, reloading`.
- **Envoy returns 503**: verify EPP health (`curl http://127.0.0.1:9090/metrics`), confirm `endpoints.yaml` lists literal IPv4 addresses (not hostnames), and check the Envoy→EPP cluster (`curl http://127.0.0.1:19000/clusters | grep ext_proc`).
- **vLLM worker unreachable**: confirm the worker is up (`curl http://<worker-ip>:8000/v1/models`) and that the address in `endpoints.yaml` matches the worker's actual IP.
- **Intel XPU device not visible inside the container**: confirm `/dev/dri` is actually passed through (`--device=/dev/dri:/dev/dri`) and the container user is in the host's `render` group — a mismatched GID here is the most common first-run failure for this substitution.

Cleanup

```
docker rm -f envoy epp vllm-0
sudo rm -rf /etc/epp /etc/envoy
```

Config Reference

Field	Value	Notes
Default model (upstream)	`Qwen/Qwen3-32B`, TP=2	same as Optimized Baseline
Discovery mechanism	file-discovery plugin, `endpoints.yaml`	literal IPv4 addresses only, no DNS
Live reload	`watchFile: true`	atomic-rename updates to `endpoints.yaml` apply without EPP restart
P/D disaggregation outside Kubernetes	combine with P/D Disaggregation's EPP plugin config, set `llm-d.ai/role: prefill`/`decode` labels per endpoint	not Intel-XPU-specific, but untested in combination here
Hardware support status	NVIDIA + vLLM only, upstream-maintained	Intel XPU is a manual reader substitution (this chapter), not a maintained path

Appendix A. Environment Variable Reference

Variable	Defined in	Purpose
REPO_ROOT	guides/env.sh	Absolute path to the repo root
GAIE_VERSION	guides/env.sh	Gateway API Inference Extension CRD version
ROUTER_CHART_VERSION	guides/env.sh	llm-d Router chart version
ROUTER_STANDALONE_CHART / ROUTER_GATEWAY_CHART	guides/env.sh	Router chart OCI addresses
NAMESPACE	user-set	Target namespace
HF_TOKEN	user-set	HuggingFace access token
GUIDE_NAME	user-set, per case	Selects the values/overlay files for a given guide

Appendix B. Known Issues

- Routing proxy init container hang: llm-d/llm-d#241 — workaround: `proxy.enabled: false`.
- vLLM usage-telemetry write failures under restricted `SecurityContext` (e.g. OpenShift): set `DO_NOT_TRACK=1`.